

Basic Graph Algorithms

Jaehyun Park

CS 97SI
Stanford University

June 29, 2015

Outline

Graphs

Adjacency Matrix and Adjacency List

Special Graphs

Depth-First and Breadth-First Search

Topological Sort

Eulerian Circuit

Minimum Spanning Tree (MST)

Strongly Connected Components (SCC)

Graphs

- ▶ An abstract way of representing connectivity using nodes (also called vertices) and edges
- ▶ We will label the nodes from 1 to n
- ▶ m edges connect some pairs of nodes
 - Edges can be either one-directional (directed) or bidirectional
- ▶ Nodes and edges can have some auxiliary information

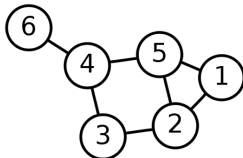


Figure from Wikipedia

Why Study Graphs?

- ▶ Lots of problems formulated and solved in terms of graphs
 - Shortest path problems
 - Network flow problems
 - Matching problems
 - 2-SAT problem
 - Graph coloring problem
 - Traveling Salesman Problem (TSP): *still unsolved!*
 - and many more...

Outline

Graphs

Adjacency Matrix and Adjacency List

Special Graphs

Depth-First and Breadth-First Search

Topological Sort

Eulerian Circuit

Minimum Spanning Tree (MST)

Strongly Connected Components (SCC)

Storing Graphs

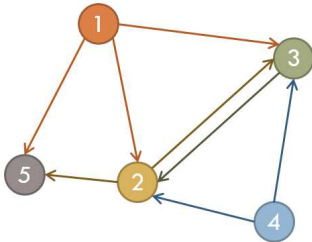
- ▶ Need to store both the set of nodes V and the set of edges E
 - Nodes can be stored in an array
 - Edges must be stored in some other way
- ▶ Want to support operations such as:
 - Retrieving all edges incident to a particular node
 - Testing if given two nodes are directly connected
- ▶ Use either adjacency matrix or adjacency list to store the edges

Adjacency Matrix

- ▶ An easy way to store connectivity information
 - Checking if two nodes are directly connected: $O(1)$ time
- ▶ Make an $n \times n$ matrix A
 - $a_{ij} = 1$ if there is an edge from i to j
 - $a_{ij} = 0$ otherwise
- ▶ Uses $\Theta(n^2)$ memory
 - Only use when n is less than a few thousands,
 - *and* when the graph is dense

Adjacency List

- ▶ Each node has a list of outgoing edges from it
 - Easy to iterate over edges incident to a certain node
 - The lists have variable lengths
 - Space usage: $\Theta(n + m)$

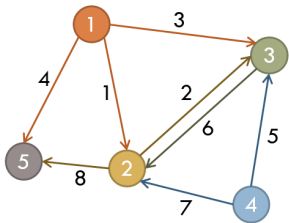


From	To		
1	2	3	5
2	3	5	
3	2		
4	2	5	
5			

Implementing Adjacency List

- ▶ Solution 1. Using linked lists
 - Too much memory/time overhead
 - Using dynamic allocated memory or pointers is bad
- ▶ Solution 2. Using an array of vectors
 - Easier to code, no bad memory issues
 - But very slow
- ▶ Solution 3. Using arrays (!)
 - Assuming the total number of edges is known
 - Very fast and memory-efficient

Implementation Using Arrays



ID	To	Next Edge ID
1	2	-
2	3	-
3	3	1
4	5	3
5	3	-
6	2	-
7	2	5
8	5	2

From	1	2	3	4	5
Last Edge ID	4	8	6	7	-

Implementation Using Arrays

- ▶ Have two arrays E of size m and LE of size n
 - E contains the edges
 - LE contains the starting pointers of the edge lists
- ▶ Initialize $LE[i] = -1$ for all i
 - $LE[i] = 0$ is also fine if the arrays are 1-indexed
- ▶ Inserting a new edge from u to v with ID k
 - $E[k].to = v$
 - $E[k].nextID = LE[u]$
 - $LE[u] = k$

Implementation Using Arrays

- ▶ Iterating over all edges starting at u:

```
for(ID = LE[u]; ID != -1; ID = E[ID].nextID)  
    // E[ID] is an edge starting from u
```

- ▶ Once built, it's hard to modify the edges
 - The graph better be static!
 - But adding more edges is easy

Outline

Graphs

Adjacency Matrix and Adjacency List

Special Graphs

Depth-First and Breadth-First Search

Topological Sort

Eulerian Circuit

Minimum Spanning Tree (MST)

Strongly Connected Components (SCC)

Tree

- ▶ A connected acyclic graph
- ▶ Most important type of special graphs
 - Many problems are easier to solve on trees
- ▶ Alternate equivalent definitions:
 - A connected graph with $n - 1$ edges
 - An acyclic graph with $n - 1$ edges
 - There is exactly one path between every pair of nodes
 - An acyclic graph but adding any edge results in a cycle
 - A connected graph but removing any edge disconnects it

Other Special Graphs

- ▶ Directed Acyclic Graph (DAG): the name says what it is
 - Equivalent to a partial ordering of nodes
- ▶ Bipartite Graph: Nodes can be separated into two groups S and T such that edges exist between S and T only (no edges within S or within T)

Outline

Graphs

Adjacency Matrix and Adjacency List

Special Graphs

Depth-First and Breadth-First Search

Topological Sort

Eulerian Circuit

Minimum Spanning Tree (MST)

Strongly Connected Components (SCC)

Graph Traversal

- ▶ The most basic graph algorithm that visits nodes of a graph in certain order
- ▶ Used as a subroutine in many other algorithms
- ▶ We will cover two algorithms
 - Depth-First Search (DFS): uses recursion (stack)
 - Breadth-First Search (BFS): uses queue

Depth-First Search

DFS(v): visits all the nodes reachable from v in depth-first order

- ▶ Mark v as visited
- ▶ For each edge $v \rightarrow u$:
 - If u is not visited, call DFS(u)
- ▶ Use non-recursive version if recursion depth is too big (over a few thousands)
 - Replace recursive calls with a stack

Breadth-First Search

$\text{BFS}(v)$: visits all the nodes reachable from v in breadth-first order

- ▶ Initialize a queue Q
- ▶ Mark v as visited and push it to Q
- ▶ While Q is not empty:
 - Take the front element of Q and call it w
 - For each edge $w \rightarrow u$:
 - ▶ If u is not visited, mark it as visited and push it to Q

Outline

Graphs

Adjacency Matrix and Adjacency List

Special Graphs

Depth-First and Breadth-First Search

Topological Sort

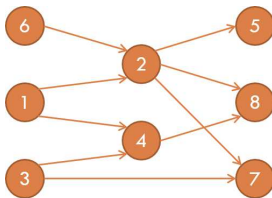
Eulerian Circuit

Minimum Spanning Tree (MST)

Strongly Connected Components (SCC)

Topological Sort

- ▶ Input: a DAG $G = (V, E)$
- ▶ Output: an ordering of nodes such that for each edge $u \rightarrow v$, u comes before v
- ▶ There can be many answers
 - e.g., both $\{6, 1, 3, 2, 7, 4, 5, 8\}$ and $\{1, 6, 2, 3, 4, 5, 7, 8\}$ are valid orderings for the graph below



Topological Sort

- ▶ Any node without an incoming edge can be the first element
 - ▶ After deciding the first node, remove outgoing edges from it
 - ▶ Repeat!
-
- ▶ Time complexity: $O(n^2 + m)$
 - Too slow...

Topological Sort (faster version)

- ▶ Precompute the number of incoming edges $\deg(v)$ for each node v
- ▶ Put all nodes v with $\deg(v) = 0$ into a queue Q
- ▶ Repeat until Q becomes empty:
 - Take v from Q
 - For each edge $v \rightarrow u$:
 - ▶ Decrement $\deg(u)$ (essentially removing the edge $v \rightarrow u$)
 - ▶ If $\deg(u) = 0$, push u to Q
- ▶ Time complexity: $\Theta(n + m)$

Outline

Graphs

Adjacency Matrix and Adjacency List

Special Graphs

Depth-First and Breadth-First Search

Topological Sort

Eulerian Circuit

Minimum Spanning Tree (MST)

Strongly Connected Components (SCC)

Eulerian Circuit

- ▶ Given an undirected graph G
- ▶ Want to find a sequence of nodes that visits every edge exactly once and comes back to the starting point

- ▶ Eulerian circuits exist if and only if
 - G is connected
 - *and* each node has an even degree

Constructive Proof of Existence

- ▶ Pick any node in G and walk randomly without using the same edge more than once
- ▶ Each node is of even degree, so when you enter a node, there will be an unused edge you exit through
 - Except at the starting point, at which you can get stuck
- ▶ When you get stuck, what you have is a cycle
 - Remove the cycle and repeat the process in each connected component
 - Glue the cycles together to finish!

Related Problems

- ▶ Eulerian path: exists if and only if the graph is connected and the number of nodes with odd degree is 0 or 2.
- ▶ Hamiltonian path/cycle: a path/cycle that visits every *node* in the graph exactly once. Looks similar but very hard (still unsolved)!

Outline

Graphs

Adjacency Matrix and Adjacency List

Special Graphs

Depth-First and Breadth-First Search

Topological Sort

Eulerian Circuit

Minimum Spanning Tree (MST)

Strongly Connected Components (SCC)

Minimum Spanning Tree (MST)

- ▶ Given an undirected weighted graph $G = (V, E)$
- ▶ Want to find a subset of E with the minimum total weight that connects all the nodes into a tree

- ▶ We will cover two algorithms:
 - Kruskal's algorithm
 - Prim's algorithm

Kruskal's Algorithm

- ▶ Main idea: the edge e^* with the smallest weight has to be in the MST
- ▶ Simple proof:
 - Assume not. Take the MST T that doesn't contain e^* .
 - Add e^* to T , which results in a cycle.
 - Remove the edge with the highest weight from the cycle.
 - ▶ The removed edge cannot be e^* since it has the smallest weight.
 - Now we have a better spanning tree than T
 - Contradiction!

Kruskal's Algorithm

- ▶ Another main idea: after an edge is chosen, the two nodes at the ends can be merged and considered as a single node (supernode)
- ▶ Pseudocode:
 - Sort the edges in increasing order of weight
 - Repeat until there is one supernode left:
 - ▶ Take the minimum weight edge e^*
 - ▶ If e^* connects two different supernodes, then connect them and merge the supernodes (use union-find)
 - Otherwise, ignore e^* and try the next edge

Prim's Algorithm

- ▶ Main idea:
 - Maintain a set S that starts out with a single node s
 - Find the smallest weighted edge $e^* = (u, v)$ that connects $u \in S$ and $v \notin S$
 - Add e^* to the MST, add v to S
 - Repeat until $S = V$
- ▶ Differs from Kruskal's in that we grow a single supernode S instead of growing multiple ones at the same time

Prim's Algorithm Pseudocode

- ▶ Initialize $S := \{s\}$, $D_v := \text{cost}(s, v)$ for every v
 - If there is no edge between s and v , $\text{cost}(s, v) = \infty$
- ▶ Repeat until $S = V$:
 - Find $v \notin S$ with smallest D_v
 - ▶ Use a priority queue or a simple linear search
 - Add v to S , add D_v to the total weight of the MST
 - For each edge (v, w) :
 - ▶ Update $D_w := \min(D_w, \text{cost}(v, w))$
- ▶ Can be modified to compute the actual MST along with the total weight

Kruskal's vs Prim's

- ▶ Kruskal's Algorithm
 - Takes $O(m \log m)$ time
 - Pretty easy to code
 - Generally slower than Prim's
- ▶ Prim's Algorithm
 - Time complexity depends on the implementation:
 - ▶ Can be $O(n^2 + m)$, $O(m \log n)$, or $O(m + n \log n)$
 - A bit trickier to code
 - Generally faster than Kruskal's

Outline

Graphs

Adjacency Matrix and Adjacency List

Special Graphs

Depth-First and Breadth-First Search

Topological Sort

Eulerian Circuit

Minimum Spanning Tree (MST)

Strongly Connected Components (SCC)

Strongly Connected Components (SCC)

- ▶ Given a *directed* graph $G = (V, E)$
- ▶ A graph is *strongly connected* if all nodes are reachable from every single node in V
- ▶ Strongly connected components of G are maximal strongly connected subgraphs of G
- ▶ The graph below has 3 SCCs: $\{a, b, e\}$, $\{c, d, h\}$, $\{f, g\}$

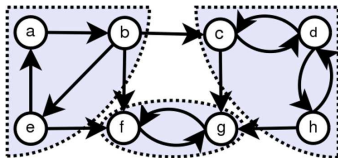


Figure from Wikipedia

Kosaraju's Algorithm

- ▶ Initialize counter $c := 0$
- ▶ While not all nodes are labeled:
 - Choose an arbitrary unlabeled node v
 - Start DFS from v
 - ▶ Check the current node x as visited
 - ▶ Recurse on all unvisited neighbors
 - ▶ After the DFS calls are finished, increment c and set the label of x as c
- ▶ Reverse the direction of all the edges
- ▶ For node v with label $n, n - 1, \dots, 1$:
 - Find all reachable nodes from v and group them as an SCC

Kosaraju's Algorithm

- ▶ We won't prove why this works
- ▶ Two graph traversals are performed
 - Running time: $\Theta(n + m)$
- ▶ Other SCC algorithms exist but this one is particularly easy to code
 - and asymptotically optimal